

CSCE 438 HW #5 (50 points)

1. *[50 points (0.65pt each)]* Choose the statements below that are **TRUE**:
 - 1.1. Strict, Sequential and Causal Consistencies are Data-Centric consistency models.
 - 1.2. Eventual Consistency, Monotonic Reads and Monotonic Writes are Client-Centric consistency models.
 - 1.3. Strict Consistency assumes absolute global time.
 - 1.4. In Strict Consistency a “read is expected to return the value resulting from the most recent write”
 - 1.5. Strict Consistency can be achieved if we have GPS.
 - 1.6. In Sequential Consistency “the result of any execution is the same as if the (read and write) operations by all processes on the data were executed in an arbitrary order.”
 - 1.7. In Causal Consistency, if an update U1 causes another update U2 to occur, then, U1 should be executed before U2 at each copy.
 - 1.8. Causal Consistency is stronger than sequential consistency.
 - 1.9. In Eventual Consistency, if no updates take place for a long time, all replicas will gradually become consistent.
 - 1.10. In a Monotonic Read, “if a process reads the value of a data item x, then any successive read operation on x by that process will always return that same value or a more recent value.”
 - 1.11. In a Write Monotonic consistent store, a write operation by a process on a data item x is completed before any successive write operation on x by the same process.
 - 1.12. One implementation of Sequential Consistency is to use a centralized process, called sequencer.
 - 1.13. Write-through will impact caching in a distributed file system.
 - 1.14. The read-ahead in distributed file systems, requests chunks of data when they are needed.
 - 1.15. The first version of NFS ran over UDP, using Sun RPC.
 - 1.16. NFS uses caching at the client (caching data, file attributes and pathname bindings).
 - 1.17. All NFS writes are write-through to disk.
 - 1.18. Inconsistencies between local caches and server in NFS are solved by comparing time-stamps.
 - 1.19. NFS always invalidated data after some time (3 seconds for data in open files).

- 1.20. NFS version 2 extends the client-side buffer caching to disk, in 64KB chunks.
- 1.21. NFS version 3 continued to use UDP, because its simplicity.
- 1.22. NFS version 3 started to support 64bit file sizes.
- 1.23. GFS was designed with Google's application workload specifically in mind.
- 1.24. In GFS, during a write operation, the master for a chunk sends data to replicas in a daisy chain.
- 1.25. In GFS, if the master reboots and then finds a chunk server has a newer version number for a chunk, it adopts that version number.
- 1.26. In GFS, if a chunk server dies, then the master decrements the count of replicas for all chunks on that chunk server.
- 1.27. GFS was designed for small streaming reads.
- 1.28. GFS was designed for large sequential writes that append.
- 1.29. In GFS, the performance of operations is very good for all apps.
- 1.30. GFS architecture contains one master server.
- 1.31. In the GFS architecture there is one chunk server for each chunk.
- 1.32. The master server in GFS holds metadata and most frequently accessed data files.
- 1.33. The master server in GFS holds all metadata in RAM.
- 1.34. A read/write operation with a chunk server in GFS specifies the chunk handle and the byte range.
- 1.35. In GFS, the client issues control/metadata requests to the chunk servers.
- 1.36. A client in GFS uses no caching.
- 1.37. The GFS consistency model defines "consistent" as: file region all clients see as same, regardless of replicas they read from."
- 1.38. In GFS, the serial success result of a write operation is DEFINED.
- 1.39. In GFS, a successful record append data mutation is DEFINED.
- 1.40. In GFS, a successful concurrent write is CONSISTENT but DEFINED.
- 1.41. In GFS, a failure of a write or record append operation is INCONSISTENT.
- 1.42. BigTable is a sparse, centralized persistent multi-dimensional sorted map.
- 1.43. The map in BigTable is indexed by row key, column key and timestamp.
- 1.44. A row range (partition) in BigTable is also called tablet.
- 1.45. The items in a BigTable cell are stored in decreasing timestamp order.
- 1.46. BigTable is an alternate way for storing data than GFS and it implements its own replication.
- 1.47. BigTable processes share the same machines with MapReduce and GFS machines.

- 1.48. In the BigTable, there is a master server and many tablet servers.
- 1.49. In the BigTable, a client communicates directly with tablet servers for reads/writes.
- 1.50. The master server in BigTable maintains the set of live tablet servers and the current assignment of tablets to tablet servers.
- 1.51. The memtable is an in RAM storage for storing the committed writes that arrived at a tablet server.
- 1.52. When memtable size increases and reaches a threshold, it is frozen and committed to an SSTable in GFS.
- 1.53. The execution of a MapReduce program creates a Master process and several Worker processes.
- 1.54. The master process of a MapReduce program assigns map and reduce tasks to worker processes.
- 1.55. The intermediate results of a MapReduce program are stored in GFS.
- 1.56. If the master process of a MapReduce program crashes a new master is elected through Chubby.
- 1.57. A Partition Function in a MapReduce program is used for ensuring that records with the same intermediate key end up at the same worker.
- 1.58. A MapReduce program will use information from GFS (location of replicas) to decide which file blocks will be processed by which worker process.
- 1.59. When a worker process in MapReduce segfaults, the information about the record it processed is lost.
- 1.60. If the master process in MapReduce sees two failures for the same record, then it tells the next worker to skip the record.
- 1.61. In HDFS, a DataNode is the same as a ChunkServer in GFS.
- 1.62. In HDFS, a NameNode is the same as the Master Node in GFS.
- 1.63. A block in GFS is the same as a chunk in HDFS.
- 1.64. If a system can provide strong consistency, the programming model (how the programmer uses the API) is greatly simplified.
- 1.65. A Mobile Cloud consists of Distributed Storage and Distributed Data Processing
- 1.66. A k-out-of-n system is an n-component system that works if and only if k or less components work
- 1.67. Radio transmission is one of the major source of energy consumption on mobile devices.
- 1.68. In “Resource allocation For edge computing” lecture data is allocated in a way that the overall data retrieval energy is minimized.

- 1.69. In “Resource allocation For edge computing” lecture, the failure probability estimation includes failures due to disconnection from network.
- 1.70. In “Resource allocation For edge computing” lecture, Importance Sampling was used for approximating the expected distance between two nodes.
- 1.71. The Read-ahead (prefetch) policy in a distributed file system, minimizes the wait when it actually is needed.
- 1.72. The Write-on-close policy in a distributed file system is synonym with session semantics.
- 1.73. The VFS layer in NFS stands for Vectored File System.
- 1.74. In MapReduce, the arguments for a mapping function is a key and a value.
- 1.75. In MapReduce, the arguments for a mapping function is a key and a value.
- 1.76. The Partition Function in MapReduce ensures that records with the same intermediate key end up at the same worker.